

✓ Heart Disease Prediction

```
# importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
sns.set()
plt.style.use('ggplot')
%matplotlib inline
```

```
#import dataset
heart_df = pd.read_csv('C:\Sunaina\Desktop\Diseas Prediction app\dataset\heart.csv')
heart_df.head(10)
```



```
# information about the dataset
heart_df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trestbps    303 non-null    int64
 4   chol        303 non-null    int64
 5   fbs         303 non-null    int64
 6   restecg     303 non-null    int64
 7   thalach     303 non-null    int64
 8   exang       303 non-null    int64
 9   oldpeak     303 non-null    float64
10   slope       303 non-null    int64
11   ca          303 non-null    int64
12   thal        303 non-null    int64
13   target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
#description about dataset
heart_df.describe()
```



```
heart_df.shape
```



```
(303, 14)
```

✓ Checking null values

```
heart_df.isnull().sum()
```



```
age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
```

```
ca      0
thal    0
target  0
dtype: int64
```

```
heart_df.notnull().sum()
```

```
age      303
sex      303
cp       303
trestbps 303
chol     303
fbs      303
restecg  303
thalach  303
exang    303
oldpeak  303
slope    303
ca       303
thal     303
target   303
dtype: int64
```

```
heart_df.dtypes
```

```
age      int64
sex      int64
cp       int64
trestbps int64
chol     int64
fbs      int64
restecg  int64
thalach  int64
exang    int64
oldpeak  float64
slope    int64
ca       int64
thal     int64
target   int64
dtype: object
```

✓ Exploratory Data Analysis(EDA)

```
#Plotting the distribution plot.
```

```
plt.figure(figsize=(20,25))
```

```
plotnumber=1
```

```
for column in heart_df:
```

```
    if plotnumber<14:
```

```
        ax=plt.subplot(4,4,plotnumber)
```

```
        sns.distplot(heart_df[column])
```

```
        plt.xlabel(column,fontsize=20)
```

```
        plt.ylabel('Values',fontsize=20)
```

```
    plotnumber+=1
```

```
plt.show()
```



```
#Correlation matrix
```

```
plt.figure(figsize = (16, 8))
```

```
corr = heart_df.corr()
```

```
mask = np.triu(np.ones_like(corr, dtype = bool))
```

```
sns.heatmap(corr, mask = mask, annot = True, fmt = '.2g', linewidths = 1)
```

```
plt.show()
```



```
#checking the variance
```

```
heart_df.var()
```

```
age      82.484558
sex      0.217166
cp       1.065132
```

```
trestbps      307.586453
chol          2686.426748
fbs           0.126877
restecg       0.276528
thalach       524.646406
exang         0.220707
oldpeak       1.348095
slope         0.379735
ca            1.045724
thal          0.374883
target        0.248836
dtype: float64
```

We can see ,there is a huge variance.So,we should normalise it.

✓ Normalization

```
heart_df['trestbps']=np.log(heart_df['trestbps'])
heart_df['chol']=np.log(heart_df['chol'])
heart_df['thalach']=np.log(heart_df['thalach'])
```

```
np.var(heart_df[["trestbps","chol","thalach"]])
```

```
↪ trestbps      0.016894
   chol         0.041401
   thalach      0.027054
   dtype: float64
```

```
heart_df.isnull().sum()
```

```
↪ age          0
   sex          0
   cp           0
   trestbps     0
   chol         0
   fbs          0
   restecg      0
   thalach      0
   exang        0
   oldpeak      0
   slope        0
   ca           0
   thal         0
   target       0
   dtype: int64
```

```
x=heart_df.drop('target',axis=1)
y=heart_df['target']
```

```
#spliting the dataset
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x, y, test_size=0.30, random_state=0)
```

```
x.info()
```

```
↪ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trestbps    303 non-null    float64
 4   chol        303 non-null    float64
 5   fbs         303 non-null    int64
 6   restecg     303 non-null    int64
 7   thalach     303 non-null    float64
 8   exang       303 non-null    int64
 9   oldpeak     303 non-null    float64
10   slope       303 non-null    int64
11   ca          303 non-null    int64
12   thal        303 non-null    int64
dtypes: float64(4), int64(9)
memory usage: 30.9 KB
```

✓ Logistic Regression

```

accuracies={}

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
lr = LogisticRegression(penalty='l2')
lr.fit(x_train,y_train)

y_pred = lr.predict(x_test)

acc=accuracy_score(y_test,y_pred)
accuracies['LR']=acc*100
print("Training accuracy score of the model is:",accuracy_score(y_train, lr.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred)*100,"%")

```

↻ Training accuracy score of the model is: 85.37735849056604 %
Testing accuracy score of the model is: 80.21978021978022 %

```

print("Confusion matrix of the model",confusion_matrix(y_test,y_pred))

print("Classification Report",classification_report(y_test,y_pred))

```

↻ Confusion matrix of the model [[32 12]
[6 41]]

Classification Report		precision	recall	f1-score	support
0	0.84	0.73	0.78	44	
1	0.77	0.87	0.82	47	
accuracy		0.80	91		
macro avg	0.81	0.80	0.80	91	
weighted avg	0.81	0.80	0.80	91	

✓ KNearestNeighbors

```

from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=8)

knn.fit(x_train,y_train)

y_pred1 = knn.predict(x_test)

acc1=accuracy_score(y_test,y_pred1)
accuracies['KNN']=acc1*100

print("Training accuracy score of the model is:",accuracy_score(y_train, knn.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred1)*100,"%")

```

↻ Training accuracy score of the model is: 85.84905660377359 %
Testing accuracy score of the model is: 75.82417582417582 %

```

print("Confusion matrix of the model",confusion_matrix(y_test,y_pred1))

print("Classification Report",classification_report(y_test,y_pred1))

```

↻ Confusion matrix of the model [[29 15]
[7 40]]

Classification Report		precision	recall	f1-score	support
0	0.81	0.66	0.72	44	
1	0.73	0.85	0.78	47	
accuracy		0.76	91		
macro avg	0.77	0.76	0.75	91	
weighted avg	0.77	0.76	0.76	91	

✓ SVM

```
from sklearn.svm import SVC

svc = SVC(probability=True)
svc.fit(x_train, y_train)

y_pred2 = svc.predict(x_test)

acc2=accuracy_score(y_test,y_pred2)
accuracies['SVM']=acc2*100

print("Training accuracy score of the model is:",accuracy_score(y_train, svc.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred2)*100,"%")
```

→ Training accuracy score of the model is: 55.660377358490564 %
Testing accuracy score of the model is: 51.64835164835166 %

```
print("Confusion matrix of the model",confusion_matrix(y_test,y_pred2))

print("Classification Report",classification_report(y_test,y_pred2))
```

→ Confusion matrix of the model [[0 44]
[0 47]]

Classification Report			precision	recall	f1-score	support
0	0.00	0.00	0.00	44		
1	0.52	1.00	0.68	47		
accuracy			0.52	91		
macro avg	0.26	0.50	0.34	91		
weighted avg	0.27	0.52	0.35	91		

✓ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier

dtc = DecisionTreeClassifier()
dtc.fit(x_train, y_train)

y_pred3 = dtc.predict(x_test)

acc3=accuracy_score(y_test,y_pred3)
accuracies['DT']=acc3*100

print("Training accuracy score of the model is:",accuracy_score(y_train, dtc.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred3)*100,"%")
```

→ Training accuracy score of the model is: 100.0 %
Testing accuracy score of the model is: 71.42857142857143 %

```
print("Confusion matrix of the model",confusion_matrix(y_test,y_pred3))

print("Classification Report",classification_report(y_test,y_pred3))
```

→ Confusion matrix of the model [[32 12]
[14 33]]

Classification Report			precision	recall	f1-score	support
0	0.70	0.73	0.71	44		
1	0.73	0.70	0.72	47		
accuracy			0.71	91		
macro avg	0.71	0.71	0.71	91		
weighted avg	0.72	0.71	0.71	91		

```
from sklearn.model_selection import GridSearchCV

grid_params = {
    'criterion' : ['gini', 'entropy'],
```

```
'max_depth' : range(2, 32, 1),
'min_samples_leaf' : range(1, 10, 1),
'min_samples_split' : range(2, 10, 1),
'splitter' : ['best', 'random']
}
```

```
grid_search = GridSearchCV(dtc, grid_params, cv = 10, n_jobs = -1, verbose = 1)
grid_search.fit(x_train, y_train)
```



```
grid_search.best_score_
```



```
np.float64(0.8538961038961039)
```

```
grid_search.best_params_
```



```
{'criterion': 'gini',
 'max_depth': 15,
 'min_samples_leaf': 7,
 'min_samples_split': 9,
 'splitter': 'random'}
```

```
dtc2 = DecisionTreeClassifier(criterion= 'entropy', max_depth= 12, min_samples_leaf= 1, min_samples_split= 2, splitter= 'random')
dtc2.fit(x_train, y_train)
```



```
y_pred4 = dtc2.predict(x_test)
acc4=accuracy_score(y_test,y_pred4)
accuracies['DT2']=acc4*100
```

```
print("Training accuracy score of the model is:",accuracy_score(y_train, dtc2.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred4)*100,"%")
```



```
Training accuracy score of the model is: 100.0 %
Testing accuracy score of the model is: 71.42857142857143 %
```

```
print("Confusion matrix of the model",confusion_matrix(y_test,y_pred4))
```

```
print("Classification Report",classification_report(y_test,y_pred4))
```



```
Confusion matrix of the model [[33 11]
 [15 32]]
```

Classification Report		precision	recall	f1-score	support
0	0.69	0.75	0.72	44	
1	0.74	0.68	0.71	47	
accuracy			0.71	91	
macro avg	0.72	0.72	0.71	91	
weighted avg	0.72	0.71	0.71	91	

```
# update dictionary
accuracies['DT']=acc4*100
del accuracies['DT2']
```

✓ Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rfc = RandomForestClassifier(criterion = 'gini', max_depth = 7, max_features = 'sqrt', min_samples_leaf = 2, min_samples_split = 4, n_estimators = 100)
rfc.fit(x_train, y_train)
```

```
y_pred5 = rfc.predict(x_test)
```

```
acc5=accuracy_score(y_test,y_pred5)
accuracies['RF']=acc5*100
```

```
print("Training accuracy score of the model is:",accuracy_score(y_train, rfc.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred5)*100,"%")
```

```

↗ Training accuracy score of the model is: 97.16981132075472 %
Testing accuracy score of the model is: 82.41758241758241 %

```

```
print("Confusion matrix of the model",confusion_matrix(y_test,y_pred5))
```

```
print("Classification Report",classification_report(y_test,y_pred5))
```

```

↗ Confusion matrix of the model [[31 13]
 [ 3 44]]
Classification Report

```

			precision	recall	f1-score	support
	0	0.91	0.70	0.79	44	
	1	0.77	0.94	0.85	47	
accuracy			0.82		91	
macro avg	0.84	0.82	0.82		91	
weighted avg	0.84	0.82	0.82		91	

✓ Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
gbc = GradientBoostingClassifier()
```

```
gbc = GradientBoostingClassifier(learning_rate = 0.05, loss = 'log_loss', n_estimators = 180)
gbc.fit(x_train, y_train)
```

```
y_pred6 = gbc.predict(x_test)
```

```
acc6 = accuracy_score(y_test,y_pred6)
accuracies['GradientBoosting']=acc6*100
```

```
print("Training accuracy score of the model is:",accuracy_score(y_train, gbc.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred6)*100,"%")
```

```

↗ Training accuracy score of the model is: 100.0 %
Testing accuracy score of the model is: 79.12087912087912 %

```

✓ XGBoost

```
from xgboost import XGBClassifier
```

```
xgb = XGBClassifier(objective = 'binary:logistic', learning_rate = 0.01, max_depth = 5, n_estimators = 180)
```

```
xgb.fit(x_train, y_train)
```

```
↗
```

```
y_pred7 = xgb.predict(x_test)
```

```
acc7=accuracy_score(y_test,y_pred7)
```

```

accuracies['XGBoost']=acc7*100
print("Training accuracy score of the model is:",accuracy_score(y_train, xgb.predict(x_train))*100,"%")
print("Testing accuracy score of the model is:",accuracy_score(y_test,y_pred7)*100,"%")

```

```

↗ Training accuracy score of the model is: 96.22641509433963 %
Testing accuracy score of the model is: 80.21978021978022 %

```

```
print("Confusion matrix of the model",confusion_matrix(y_test,y_pred7))
```

```
print("Classification Report",classification_report(y_test,y_pred7))
```

```

↗ Confusion matrix of the model [[32 12]
 [ 6 41]]
Classification Report

```

			precision	recall	f1-score	support
	0	0.84	0.73	0.78	44	
	1	0.77	0.87	0.82	47	

accuracy			0.80	91
macro avg	0.81	0.80	0.80	91
weighted avg	0.81	0.80	0.80	91

```
colors = ["purple", "green", "orange", "magenta", "blue", "black"]
```

```
# sns.set_style("whitegrid")
plt.figure(figsize=(16,8))
plt.xticks(np.arange(0,1200,10))
plt.ylabel("Accuracy %")
plt.xlabel("Algorithms")
sns.barplot(x=list(accuracies.keys()), y=list(accuracies.values()), palette=colors )
plt.show()
```



```
models = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'SVM', 'Decision Tree', 'Random Forest', 'Gradient Boosting', 'XgBoost'],
    'Score': [acc, acc1, acc2, acc4, acc5, acc6, acc7]
})
```

```
models.sort_values(by = 'Score', ascending = False)
```



```
import pickle
model = rfc
pickle.dump(model, open("heart.pkl", 'wb'))
```

```
from sklearn import metrics
plt.figure(figsize=(8,5))
models = [
    {
        'label': 'LR',
        'model': lr,
    },
    {
        'label': 'DT',
        'model': dtc2,
    },
    {
        'label': 'SVM',
        'model': svc,
    },
    {
        'label': 'KNN',
        'model': knn,
    },
    {
        'label': 'XGBoost',
        'model': xgb,
    },
    {
        'label': 'RF',
        'model': rfc,
    },
    {
        'label': 'GBDT',
        'model': gbc,
    }
]
for m in models:
    model = m['model']
    model.fit(x_train, y_train)
    y_pred=model.predict(x_test)
    fpr1, tpr1, thresholds = metrics.roc_curve(y_test, model.predict_proba(x_test)[:,-1])
    auc = metrics.roc_auc_score(y_test,model.predict(x_test))
    plt.plot(fpr1, tpr1, label='%s - ROC (area = %0.2f)' % (m['label'], auc))

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1 - Specificity (False Positive Rate)', fontsize=12)
plt.ylabel('Sensitivity (True Positive Rate)', fontsize=12)
```



```
plt.title('ROC - Heart Disease Prediction', fontsize=12)
plt.legend(loc="lower right", fontsize=12)
plt.savefig("roc_heart.jpeg", format='jpeg', dpi=400, bbox_inches='tight')
plt.show()
```



```
from sklearn import metrics
import numpy as np
import matplotlib.pyplot as plt
models = [
{
    'label': 'LR',
    'model': lr,
},
{
    'label': 'DT',
    'model': dtc2,
},
{
    'label': 'SVM',
    'model': svc,
},
{
    'label': 'KNN',
    'model': knn,
},
{
    'label': 'XGBoost',
    'model': xgb,
},
{
    'label': 'RF',
    'model': rfc,
},
{
    'label': 'GBDT',
    'model': gbc,
}
]

means_roc = []
means_accuracy = [100*round(acc,4), 100*round(acc4,4), 100*round(acc2,4), 100*round(acc1,4), 100*round(acc7,4),
                  100*round(acc5,4), 100*round(acc6,4)]

for m in models:
    model = m['model']
    model.fit(x_train, y_train)
    y_pred=model.predict(x_test)
    fpr1, tpr1, thresholds = metrics.roc_curve(y_test, model.predict_proba(x_test)[: ,1])
    auc = metrics.roc_auc_score(y_test,model.predict(x_test))
    auc = 100*round(auc,4)
    means_roc.append(auc)

print(means_accuracy)
print(means_roc)

# data to plot
n_groups = 7
means_accuracy = tuple(means_accuracy)
means_roc = tuple(means_roc)

# create plot
fig, ax = plt.subplots(figsize=(8,5))
index = np.arange(n_groups)
bar_width = 0.35
opacity = 0.8

rects1 = plt.bar(index, means_accuracy, bar_width,
alpha=opacity,
color='mediumpurple',
label='Accuracy (%)')

rects2 = plt.bar(index + bar_width, means_roc, bar_width,
alpha=opacity,
color='rebeccapurple',
label='ROC (%)')
```

```
plt.xlim([-1, 8])
plt.ylim([70, 105])

plt.title('Performance Evaluation - Heart Disease Prediction', fontsize=12)
plt.xticks(index, (' LR', ' DT', ' SVM', ' KNN', 'XGBoost', ' RF', ' GBDT'), rotation=40, ha='center', fontsize=12)
plt.legend(loc="upper right", fontsize=10)
plt.savefig("PE_heart.jpeg", format='jpeg', dpi=400, bbox_inches='tight')
plt.show()
```

